

```
1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4 import seaborn as sns

1 df = pd.read_csv('https://github.com/YBI-Foundation/Dataset/raw/main/Bank%20Churn%20Modelling.csv')
2 print(df)
```

	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	\
0	15634602	Hargrave	619	France	Female	42	2	
1	15647311	Hill	608	Spain	Female	41	1	
2	15619304	Onio	502	France	Female	42	8	
3	15701354	Boni	699	France	Female	39	1	
4	15737888	Mitchell	850	Spain	Female	43	2	
...	
9995	15606229	Obijiaku	771	France	Male	39	5	
9996	15569892	Johnstone	516	France	Male	35	10	
9997	15584532	Liu	709	France	Female	36	7	
9998	15682355	Sabbatini	772	Germany	Male	42	3	
9999	15628319	Walker	792	France	Female	28	4	

	Balance	Num Of Products	Has Credit Card	Is Active Member	\
0	0.00	1	1	1	
1	83807.86	1	0	1	
2	159660.80	3	1	0	
3	0.00	2	0	0	
4	125510.82	1	1	1	
...	
9995	0.00	2	1	0	
9996	57369.61	1	1	1	
9997	0.00	1	0	1	
9998	75075.31	2	1	0	
9999	130142.79	1	1	0	

	Estimated Salary	Churn
0	101348.88	1
1	112542.58	0
2	113931.57	1
3	93826.63	0
4	79084.10	0
...
9995	96270.64	0
9996	101699.77	0
9997	42085.58	1
9998	92888.52	1
9999	38190.78	0

[10000 rows x 13 columns]

```
1 df.head()
```

	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	Num Of Products	Has Credit Card	Is Active Member	E
0	15634602	Hargrave	619	France	Female	42	2	0.00	1	1	1	
1	15647311	Hill	608	Spain	Female	41	1	83807.86	1	0	1	
2	15619304	Onio	502	France	Female	42	8	159660.80	3	1	0	
3	15701354	Boni	699	France	Female	39	1	0.00	2	0	0	
4	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	1	1	

```
1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   CustomerId            10000 non-null  int64
1   Surname               10000 non-null  object
2   CreditScore           10000 non-null  int64
3   Geography             10000 non-null  object
4   Gender                10000 non-null  object
5   Age                   10000 non-null  int64
6   Tenure                10000 non-null  int64
7   Balance               10000 non-null  float64
8   Num Of Products       10000 non-null  int64
9   Has Credit Card       10000 non-null  int64
10  Is Active Member      10000 non-null  int64
11  Estimated Salary      10000 non-null  float64
12  Churn                 10000 non-null  int64
dtypes: float64(2), int64(8), object(3)
memory usage: 1015.8+ KB
```

```

1 df.duplicated('CustomerId').sum()

0

1 df = df.set_index('CustomerId')

1 df.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 10000 entries, 15634602 to 15628319
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Surname                10000 non-null  object
1   CreditScore             10000 non-null  int64
2   Geography               10000 non-null  object
3   Gender                  10000 non-null  object
4   Age                     10000 non-null  int64
5   Tenure                  10000 non-null  int64
6   Balance                 10000 non-null  float64
7   Num Of Products         10000 non-null  int64
8   Has Credit Card         10000 non-null  int64
9   Is Active Member        10000 non-null  int64
10  Estimated Salary        10000 non-null  float64
11  Churn                    10000 non-null  int64
dtypes: float64(2), int64(7), object(3)
memory usage: 1015.6+ KB

```

```
1 df.describe()
```

	CreditScore	Age	Tenure	Balance	Num Of Products	Has Credit Card	Is Active Member	Estimated Salary
count	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000
mean	650.528800	38.921800	5.012800	76485.889288	1.530200	0.70550	0.515100	100090.239881
std	96.653299	10.487806	2.892174	62397.405202	0.581654	0.45584	0.499797	57510.492818
min	350.000000	18.000000	0.000000	0.000000	1.000000	0.00000	0.000000	11.580000
25%	584.000000	32.000000	3.000000	0.000000	1.000000	0.00000	0.000000	51002.110000
50%	652.000000	37.000000	5.000000	97198.540000	1.000000	1.00000	1.000000	100193.915000
75%	718.000000	44.000000	7.000000	127644.240000	2.000000	1.00000	1.000000	149388.247500
max	850.000000	92.000000	10.000000	250898.090000	4.000000	1.00000	1.000000	199992.480000

Encoding

```

1 df['Geography'].value_counts()

France    5014
Germany   2509
Spain     2477
Name: Geography, dtype: int64

1 df.replace({'Geography':{'France':2,'Germany':1,'Spain':0}},inplace=True)

1 df['Gender'].value_counts()

Male      5457
Female    4543
Name: Gender, dtype: int64

1 df.replace({'Gender':{'Male':0,'Female':1}},inplace=True)

1 df['Num Of Products'].value_counts()

1     5084
2     4590
3      266
4       60
Name: Num Of Products, dtype: int64

1 df.replace({'Num Of Products':{'1:0,2:1,3:1,4:1}},inplace=True)

```

```
1 df['Has Credit Card'].value_counts()
```

```
1    7055
0    2945
Name: Has Credit Card, dtype: int64
```

```
1 df['Is Active Member'].value_counts()
```

```
1    5151
0    4849
Name: Is Active Member, dtype: int64
```


```
1 df.loc[(df['Balance'] == 0), 'Churn'].value_counts()
```

```
0    3117
1     500
Name: Churn, dtype: int64
```

```
1 df['Zer Balance'] = np.where(df['Balance']>0,1,0)
```

```
1 df['Zer Balance'].hist()
```

<Axes: >



```
1 df.groupby(['Churn', 'Geography']).count()
```

	Surname	CreditScore	Gender	Age	Tenure	Balance	Num Of Products	Has Credit Card	Is Active Member	Estimated
Churn	Geography									
0	0	2064	2064	2064	2064	2064	2064	2064	2064	2064
	1	1695	1695	1695	1695	1695	1695	1695	1695	1695
	2	4204	4204	4204	4204	4204	4204	4204	4204	4204
1	0	413	413	413	413	413	413	413	413	413
	1	814	814	814	814	814	814	814	814	814
	2	810	810	810	810	810	810	810	810	810

Define Label and Features

```
1 df.columns
```

```
Index(['Surname', 'CreditScore', 'Geography', 'Gender', 'Age', 'Tenure',
      'Balance', 'Num Of Products', 'Has Credit Card', 'Is Active Member',
      'Estimated Salary', 'Churn', 'Zer Balance'],
      dtype='object')
```


```
1 x = df.drop(['Surname', 'Churn'], axis = 1)
```

```
1 y = df['Churn']
```

```
1 df['Churn'].value_counts()
```

```
0    7963
1    2037
Name: Churn, dtype: int64
```

```
1 sns.countplot(x = 'Churn', data = df);
```



```
1 x.shape,y.shape
```

```
((10000, 11), (10000,))
```

Random Under Sampling

```
1 from imblearn.under_sampling import RandomUnderSampler
```

```
1 rus = RandomUnderSampler(random_state=2529)
```

```
1 x_rus,y_rus = rus.fit_resample(x,y)
```

```
1 x_rus.shape,y_rus.shape,x.shape,y.shape
```

```
((4074, 11), (4074,), (10000, 11), (10000,))
```

```
1 y.value_counts()
```

```
0    7963
1    2037
Name: Churn, dtype: int64
```

```
1 y_rus.value_counts()
```

```
0    2037
1    2037
Name: Churn, dtype: int64
```

```
1 y_rus.plot(kind = 'hist')
```

<Axes: ylabel='Frequency'>

**Random Over Sampling**

```

1 from imblearn.over_sampling import RandomOverSampler
1 ros = RandomOverSampler(random_state=2529)
1 x_ros,y_ros = ros.fit_resample(x,y)
1 x_ros.shape,y_ros.shape,x.shape,y.shape

((15926, 11), (15926,)), (10000, 11), (10000,))

```

```

1 y.value_counts()

0    7963
1    2037
Name: Churn, dtype: int64

```

```

1 y_ros.value_counts()

1    7963
0    7963
Name: Churn, dtype: int64


```

```

1 y_ros.plot(kind = 'hist')

```

<Axes: ylabel='Frequency'>

**Train Test Split**

```

1 from sklearn.model_selection import train_test_split

```

Split Original Data

```

1 x_train,x_test,y_train,y_test = train_test_split(x,y,test_size = 0.3,random_state=2529)

```

Split Random Under Sample Data

```

1 x_train_rus,x_test_rus,y_train_rus,y_test_rus = train_test_split(x_rus,y_rus,test_size = 0.3,random_state=2529)

```

Split Random Over Sample Data

```
1 x_train_ros,x_test_ros,y_train_ros,y_test_ros = train_test_split(x_ros,y_ros,test_size = 0.3,random_state=2529)
```

Standard Features

```
1 from sklearn.preprocessing import StandardScaler
2 sc = StandardScaler()
```

Standard Original Data

```
1 x_train[['CreditScore','Age','Tenure','Balance','Estimated Salary']] = sc.fit_transform(x_train[['CreditScore','Age','Tenure','
1 x_test[['CreditScore','Age','Tenure','Balance','Estimated Salary']] = sc.fit_transform(x_test[['CreditScore','Age','Tenure','
```

Standardize Random Under Sample Data

```
1 x_train_rus[['CreditScore','Age','Tenure','Balance','Estimated Salary']] = sc.fit_transform(x_train_rus[['CreditScore','Age','T
1 x_test_rus[['CreditScore','Age','Tenure','Balance','Estimated Salary']] = sc.fit_transform(x_test_rus[['CreditScore','Age','T
```

Standardize Random over Sample Data

```
1 x_train_ros[['CreditScore','Age','Tenure','Balance','Estimated Salary']] = sc.fit_transform(x_train_ros[['CreditScore','Age','
1 x_test_ros[['CreditScore','Age','Tenure','Balance','Estimated Salary']] = sc.fit_transform(x_test_ros[['CreditScore','Age','T
```

Support Vector Machine Classifier

```
1 from sklearn.svm import SVC
2 svc = SVC()
```

```
1 svc.fit(x_train,y_train)
```

▼ SVC
SVC()

```
1 y_pred = svc.predict(x_test)
```

Model Accuracy

```
1 from sklearn.metrics import confusion_matrix,classification_report
```

```
1 confusion_matrix(y_test,y_pred)
```

```
array([[2381,  33],
       [ 436, 150]])
```

```
1 print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
0	0.85	0.99	0.91	2414
1	0.82	0.26	0.39	586
accuracy			0.84	3000
macro avg	0.83	0.62	0.65	3000
weighted avg	0.84	0.84	0.81	3000

Hyperparameter Tunning

```
1 from sklearn.model_selection import GridSearchCV
```

```

1 param_grid = {'C':[0.1,1,10],
2               'gamma':[1,0.1,0.01],
3               'kernel':['rbf'],
4               'class_weight':['balanced']}

```

```

1 grid = GridSearchCV(SVC(),param_grid,refit = True,verbose = 2, cv = 2)
2 grid.fit(x_train,y_train)

```

```

Fitting 2 folds for each of 9 candidates, totalling 18 fits
[CV] END ..C=0.1, class_weight=balanced, gamma=1, kernel=rbf; total time= 2.6s
[CV] END ..C=0.1, class_weight=balanced, gamma=1, kernel=rbf; total time= 3.0s
[CV] END C=0.1, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 1.3s
[CV] END C=0.1, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 1.1s
[CV] END C=0.1, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 1.2s
[CV] END C=0.1, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 1.2s
[CV] END ...C=1, class_weight=balanced, gamma=1, kernel=rbf; total time= 1.3s
[CV] END ...C=1, class_weight=balanced, gamma=1, kernel=rbf; total time= 1.3s
[CV] END ..C=1, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 0.9s
[CV] END ..C=1, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 1.0s
[CV] END ..C=1, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 1.0s
[CV] END .C=1, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 1.9s
[CV] END ...C=10, class_weight=balanced, gamma=1, kernel=rbf; total time= 1.8s
[CV] END ...C=10, class_weight=balanced, gamma=1, kernel=rbf; total time= 1.3s
[CV] END .C=10, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 1.1s
[CV] END .C=10, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 1.1s
[CV] END C=10, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 1.0s
[CV] END C=10, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 1.0s

```

```

> GridSearchCV
> estimator: SVC
  > SVC

```

```
1 print(grid.best_estimator_)
```

```
SVC(C=10, class_weight='balanced', gamma=1)
```

```
1 grid_predictions = grid.predict(x_test)
```

```
1 confusion_matrix(y_test,grid_predictions)
```

```

array([[2159, 255],
       [ 343, 243]])

```

```
1 print(classification_report(y_test,grid_predictions))
```

	precision	recall	f1-score	support
0	0.86	0.89	0.88	2414
1	0.49	0.41	0.45	586
accuracy			0.80	3000
macro avg	0.68	0.65	0.66	3000
weighted avg	0.79	0.80	0.79	3000

Model with Random Under Sampling

```
1 svc_rus = SVC()
```

```
1 svc_rus.fit(x_train_rus,y_train_rus)
```

```

> SVC
SVC()

```

```
1 y_pred_rus = svc_rus.predict(x_test_rus)
```

Model Accuracy

```
1 confusion_matrix(y_test_rus,y_pred_rus)
```

```

array([[470, 157],
       [174, 422]])

```

```
1 print(classification_report(y_test_rus,y_pred_rus))
```

	precision	recall	f1-score	support
0	0.73	0.75	0.74	627
1	0.73	0.71	0.72	596
accuracy			0.73	1223
macro avg	0.73	0.73	0.73	1223
weighted avg	0.73	0.73	0.73	1223

Hyperparameter Tunning

```
1 param_grid = {'C':[0.1,1,10],
2               'gamma':[1,0.1,0.01],
3               'kernel':['rbf'],
4               'class_weight':['balanced']}
```

```
1 grid_rus = GridSearchCV(SVC(),param_grid,refit = True,verbose = 2, cv = 2)
2 grid_rus.fit(x_train_rus,y_train_rus)
```

```
Fitting 2 folds for each of 9 candidates, totalling 18 fits
[CV] END ..C=0.1, class_weight=balanced, gamma=1, kernel=rbf; total time= 1.2s
[CV] END ..C=0.1, class_weight=balanced, gamma=1, kernel=rbf; total time= 2.4s
[CV] END C=0.1, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 0.4s
[CV] END C=0.1, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 0.4s
[CV] END C=0.1, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 0.5s
[CV] END C=0.1, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 0.4s
[CV] END ...C=1, class_weight=balanced, gamma=1, kernel=rbf; total time= 0.4s
[CV] END ...C=1, class_weight=balanced, gamma=1, kernel=rbf; total time= 0.4s
[CV] END ..C=1, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 0.3s
[CV] END ..C=1, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 0.2s
[CV] END .C=1, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 0.2s
[CV] END .C=1, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 0.2s
[CV] END ...C=10, class_weight=balanced, gamma=1, kernel=rbf; total time= 0.2s
[CV] END ...C=10, class_weight=balanced, gamma=1, kernel=rbf; total time= 0.2s
[CV] END .C=10, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 0.2s
[CV] END .C=10, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 0.2s
[CV] END C=10, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 0.2s
[CV] END C=10, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 0.2s
```

```

> GridSearchCV
> estimator: SVC
  > SVC

```

```
1 print(grid_rus.best_estimator_)
```

```
SVC(C=1, class_weight='balanced', gamma=0.1)
```

```
1 grid_predictions_rus = grid_rus.predict(x_test_rus)
```

```
1 confusion_matrix(y_test_rus,grid_predictions_rus)
```

```
array([[476, 151],
       [172, 424]])
```

```
1 print(classification_report(y_test_rus,grid_predictions_rus))
```

	precision	recall	f1-score	support
0	0.73	0.76	0.75	627
1	0.74	0.71	0.72	596
accuracy			0.74	1223
macro avg	0.74	0.74	0.74	1223
weighted avg	0.74	0.74	0.74	1223

Model with Random Over Sampling

```
1 svc_ros = SVC()
```

```
1 svc_ros.fit(x_train_ros,y_train_ros)
```




```
1 y_pred_ros = svc_ros.predict(x_test_ros)
```

Model Accuracy

```
1 confusion_matrix(y_test_ros,y_pred_ros)
```

```
array([[1823,  556],
       [ 626, 1773]])
```

```
1 print(classification_report(y_test_ros,y_pred_ros))
```


	precision	recall	f1-score	support
0	0.74	0.77	0.76	2379
1	0.76	0.74	0.75	2399
accuracy			0.75	4778
macro avg	0.75	0.75	0.75	4778
weighted avg	0.75	0.75	0.75	4778

Hyperparameter Tunning

```
1 param_grid = {'C':[0.1,1,10],
2               'gamma':[1,0.1,0.01],
3               'kernel':['rbf'],
4               'class_weight':['balanced']}
```

```
1 grid_ros = GridSearchCV(SVC(),param_grid,refit = True,verbose = 2, cv = 2)
2 grid_ros.fit(x_train_ros,y_train_ros)
```

```
Fitting 2 folds for each of 9 candidates, totalling 18 fits
[CV] END ..C=0.1, class_weight=balanced, gamma=1, kernel=rbf; total time= 6.7s
[CV] END ..C=0.1, class_weight=balanced, gamma=1, kernel=rbf; total time= 3.8s
[CV] END C=0.1, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 2.8s
[CV] END C=0.1, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 3.5s
[CV] END C=0.1, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 3.7s
[CV] END C=0.1, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 3.0s
[CV] END ...C=1, class_weight=balanced, gamma=1, kernel=rbf; total time= 3.2s
[CV] END ...C=1, class_weight=balanced, gamma=1, kernel=rbf; total time= 4.7s
[CV] END ..C=1, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 2.4s
[CV] END ..C=1, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 2.4s
[CV] END .C=1, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 2.7s
[CV] END .C=1, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 5.1s
[CV] END ...C=10, class_weight=balanced, gamma=1, kernel=rbf; total time= 3.1s
[CV] END ...C=10, class_weight=balanced, gamma=1, kernel=rbf; total time= 2.9s
[CV] END .C=10, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 2.8s
[CV] END .C=10, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 4.0s
[CV] END C=10, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 2.9s
[CV] END C=10, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 2.5s
```



```
1 print(grid_ros.best_estimator_)
```

```
SVC(C=10, class_weight='balanced', gamma=1)
```

```
1 grid_predictions_ros = grid_ros.predict(x_test_ros)
```

```
1 confusion_matrix(y_test_ros,grid_predictions_ros)
```

```
array([[2047,  332],
       [  68, 2331]])
```

```
1 print(classification_report(y_test_ros,grid_predictions_ros))
```

	precision	recall	f1-score	support
0	0.97	0.86	0.91	2379
1	0.88	0.97	0.92	2399

accuracy			0.92	4778
macro avg	0.92	0.92	0.92	4778
weighted avg	0.92	0.92	0.92	4778

Let's Compare

```
1 print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
0	0.85	0.99	0.91	2414
1	0.82	0.26	0.39	586
accuracy			0.84	3000
macro avg	0.83	0.62	0.65	3000
weighted avg	0.84	0.84	0.81	3000

```
1 print(classification_report(y_test,grid_predictions))
```

	precision	recall	f1-score	support
0	0.86	0.89	0.88	2414
1	0.49	0.41	0.45	586
accuracy			0.80	3000
macro avg	0.68	0.65	0.66	3000
weighted avg	0.79	0.80	0.79	3000

```
1 print(classification_report(y_test_rus,y_pred_rus))
```

	precision	recall	f1-score	support
0	0.73	0.75	0.74	627
1	0.73	0.71	0.72	596
accuracy			0.73	1223
macro avg	0.73	0.73	0.73	1223
weighted avg	0.73	0.73	0.73	1223

```
1 print(classification_report(y_test_rus,grid_predictions_rus))
```

	precision	recall	f1-score	support
0	0.73	0.76	0.75	627
1	0.74	0.71	0.72	596
accuracy			0.74	1223
macro avg	0.74	0.74	0.74	1223
weighted avg	0.74	0.74	0.74	1223

```
1 print(classification_report(y_test_ros,y_pred_ros))
```

	precision	recall	f1-score	support
0	0.74	0.77	0.76	2379
1	0.76	0.74	0.75	2399
accuracy			0.75	4778
macro avg	0.75	0.75	0.75	4778
weighted avg	0.75	0.75	0.75	4778

```
1 print(classification_report(y_test_ros,grid_predictions_ros))
```

	precision	recall	f1-score	support
0	0.97	0.86	0.91	2379
1	0.88	0.97	0.92	2399
accuracy			0.92	4778
macro avg	0.92	0.92	0.92	4778
weighted avg	0.92	0.92	0.92	4778

